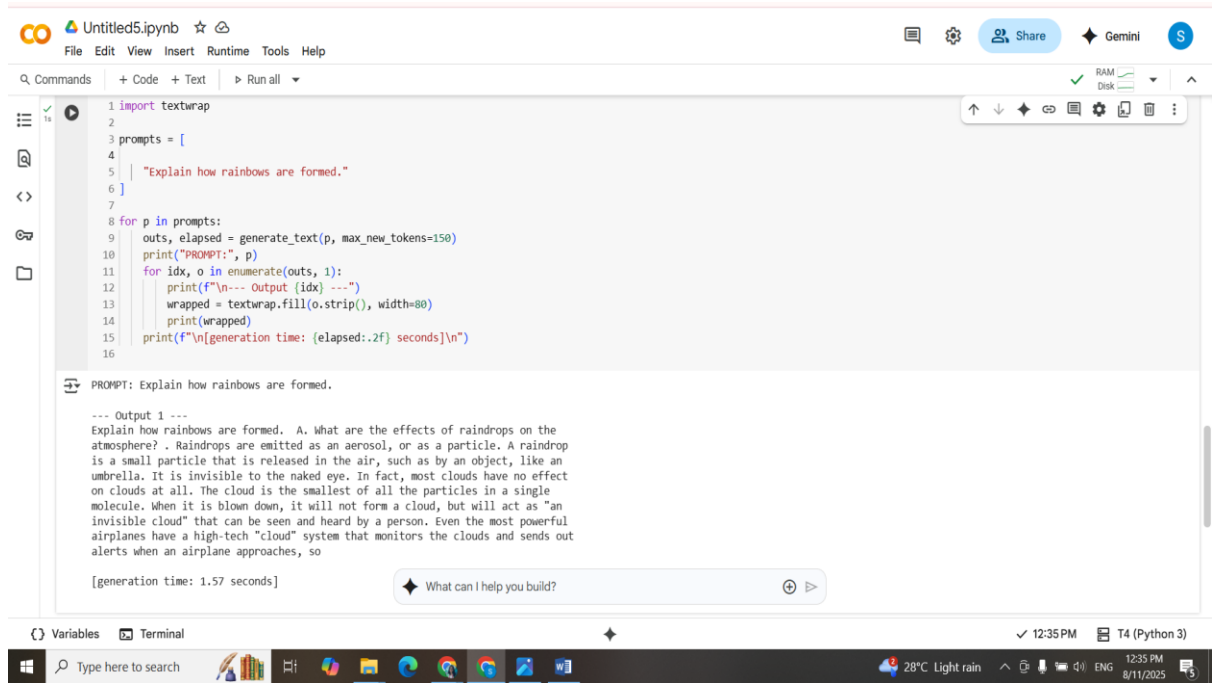


Assignment from Session 3 and session 4 (Basics of Gen AI and Basics of RAG)

1. Write a Python script to send a prompt "Explain how rainbows are formed" using OpenAI's GPT-3.5 Turbo or Hugging Face Transformers and print the response.



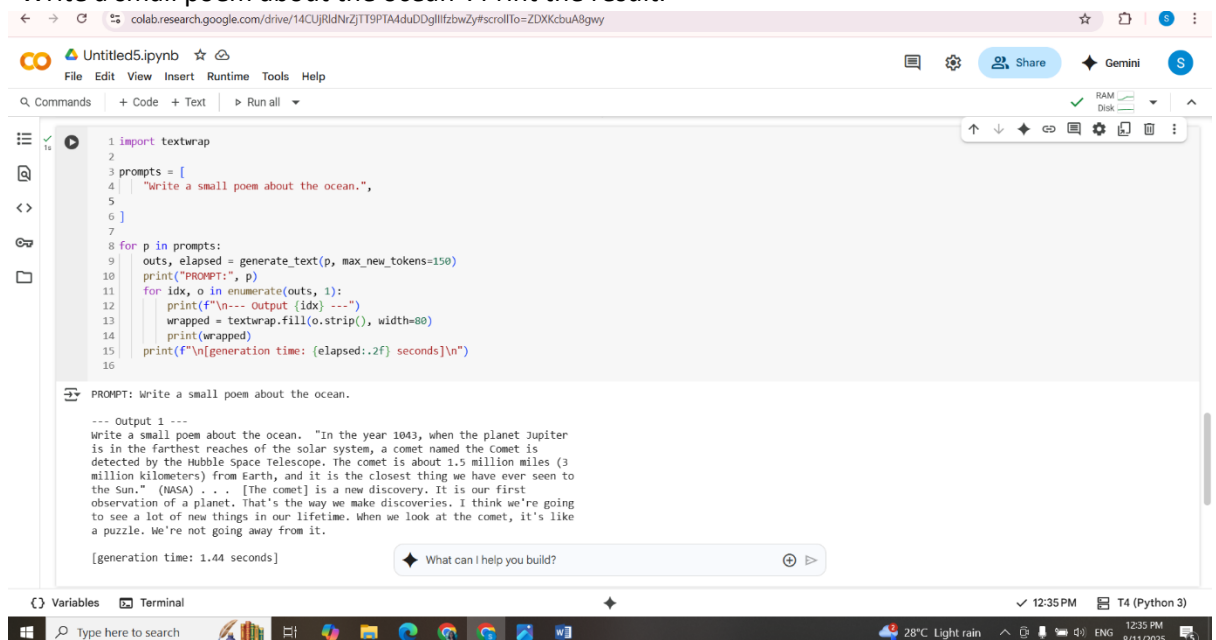
```
1 import textwrap
2
3 prompts = [
4     "Explain how rainbows are formed."
5 ]
6
7
8 for p in prompts:
9     outs, elapsed = generate_text(p, max_new_tokens=150)
10    print(f"PROMPT: {p}")
11    for idx, o in enumerate(outs, 1):
12        print(f"--- Output {idx} ---")
13        wrapped = textwrap.fill(o.strip(), width=80)
14        print(wrapped)
15    print(f"\n[generation time: {elapsed:.2f} seconds]\n")
16
```

PROMPT: Explain how rainbows are formed.

--- Output 1 ---
Explain how rainbows are formed. A. What are the effects of raindrops on the atmosphere? . Raindrops are emitted as an aerosol, or as a particle. A raindrop is a small particle that is released in the air, such as by an object, like an umbrella. It is invisible to the naked eye. In fact, most clouds have no effect on clouds at all. The cloud is the smallest of all the particles in a single molecule. When it is blown down, it will not form a cloud, but will act as "an invisible cloud" that can be seen and heard by a person. Even the most powerful airplanes have a high-tech "cloud" system that monitors the clouds and sends out alerts when an airplane approaches, so

[generation time: 1.57 seconds]

2. Use the transformers library from Hugging Face to load a model and generate text from a prompt:
"Write a small poem about the ocean". Print the result.



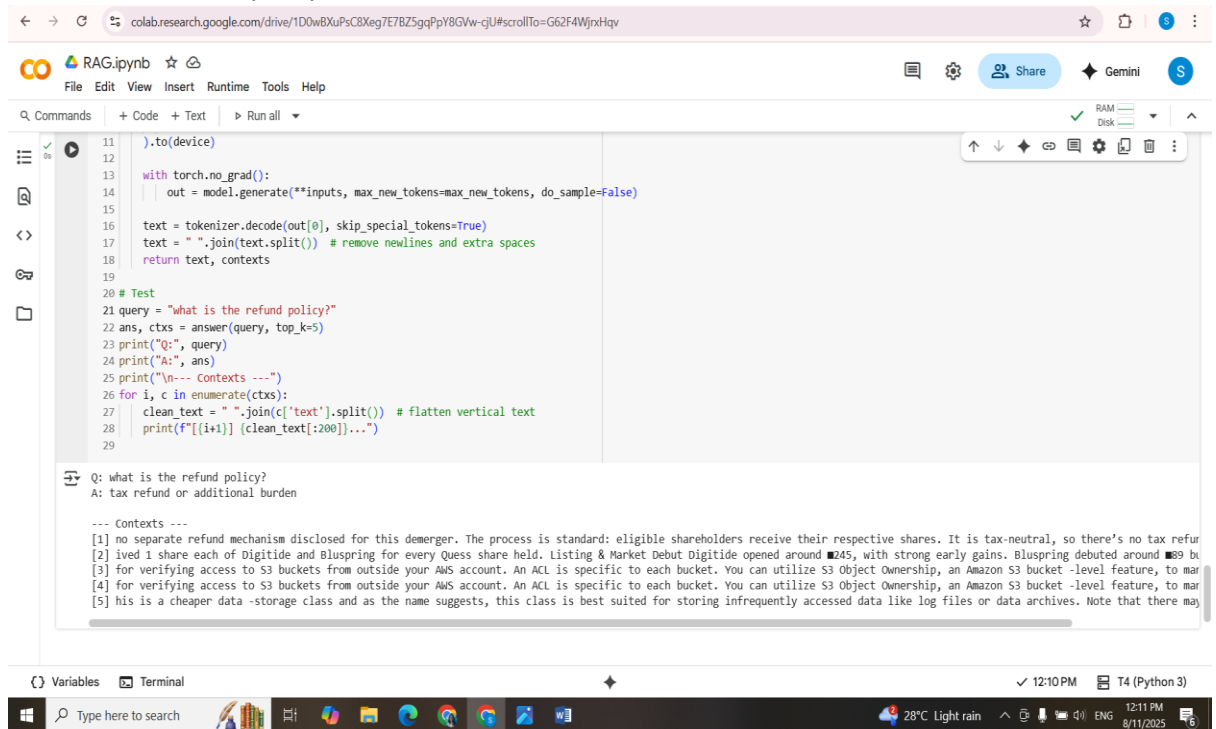
```
1 import textwrap
2
3 prompts = [
4     "Write a small poem about the ocean."
5 ]
6
7
8 for p in prompts:
9     outs, elapsed = generate_text(p, max_new_tokens=150)
10    print(f"PROMPT: {p}")
11    for idx, o in enumerate(outs, 1):
12        print(f"--- Output {idx} ---")
13        wrapped = textwrap.fill(o.strip(), width=80)
14        print(wrapped)
15    print(f"\n[generation time: {elapsed:.2f} seconds]\n")
16
```

PROMPT: Write a small poem about the ocean.

--- Output 1 ---
Write a small poem about the ocean. "In the year 1043, when the planet Jupiter is in the farthest reaches of the solar system, a comet named the Comet is detected by the Hubble Space Telescope. The comet is about 1.5 million miles (3 million kilometers) from Earth, and it is the closest thing we have ever seen to the Sun." (NASA) . . . [The comet] is a new discovery. It is our first observation of a planet. That's the way we make discoveries. I think we're going to see a lot of new things in our lifetime. When we look at the comet, it's like a puzzle. We're not going away from it.

[generation time: 1.44 seconds]

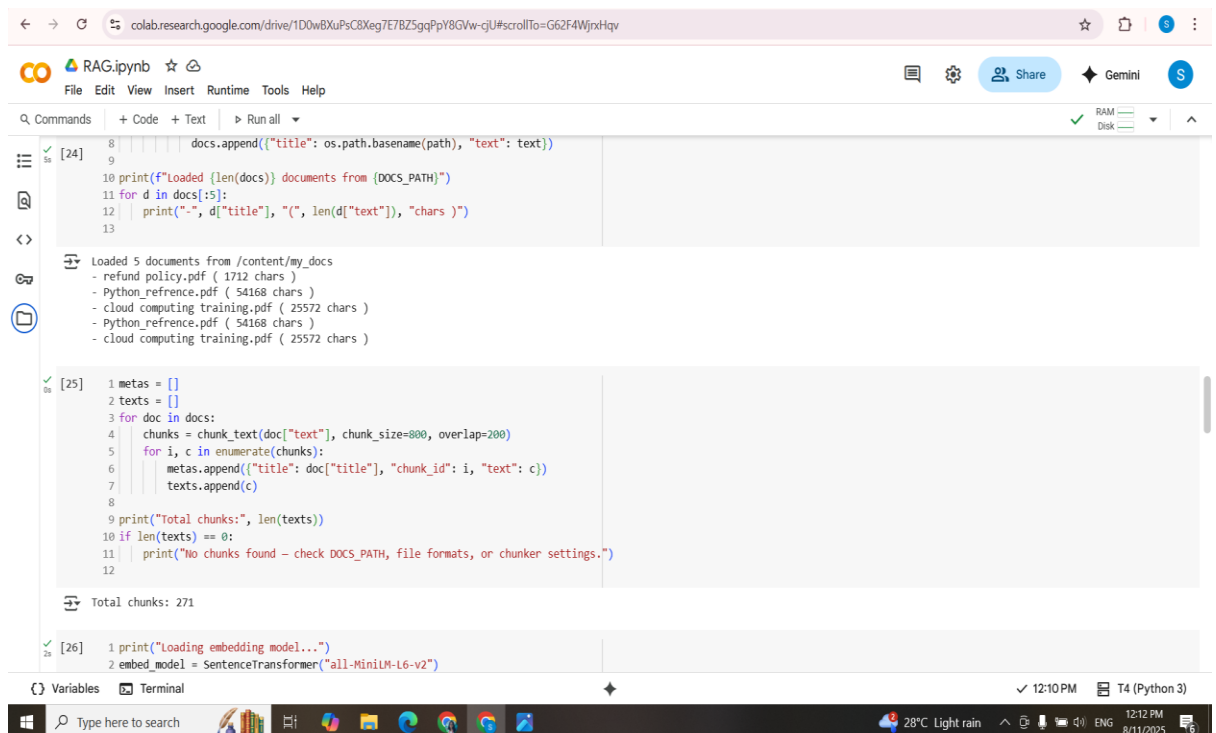
3. Using RetrievalQA, create a pipeline that can answer the question:
“What is the refund policy?”



```
11 ).to(device)
12
13 with torch.no_grad():
14     out = model.generate(**inputs, max_new_tokens=max_new_tokens, do_sample=False)
15
16 text = tokenizer.decode(out[0], skip_special_tokens=True)
17 text = " ".join(text.split()) # remove newlines and extra spaces
18 return text, contexts
19
20 # Test
21 query = "what is the refund policy?"
22 ans, ctxs = answer(query, top_k=5)
23 print("Q:", query)
24 print("A:", ans)
25 print("\n--- Contexts ---")
26 for i, c in enumerate(ctxs):
27     clean_text = " ".join(c['text'].split()) # flatten vertical text
28     print(f"[{i+1}] {clean_text[:200]}...")
29
Q: what is the refund policy?
A: tax refund or additional burden

--- Contexts ---
[1] no separate refund mechanism disclosed for this demerger. The process is standard: eligible shareholders receive their respective shares. It is tax-neutral, so there's no tax refu
[2] ived 1 share each of Digitide and Bluspring for every Qness share held. Listing & Market Debut Digitide opened around 2024, with strong early gains. Bluspring debuted around 2019 bu
[3] for verifying access to S3 buckets from outside your AWS account. An ACL is specific to each bucket. You can utilize S3 Object Ownership, an Amazon S3 bucket -level feature, to mar
[4] for verifying access to S3 buckets from outside your AWS account. An ACL is specific to each bucket. You can utilize S3 Object Ownership, an Amazon S3 bucket -level feature, to mar
[5] his is a cheaper data -storage class and as the name suggests, this class is best suited for storing infrequently accessed data like log files or data archives. Note that there may
```

4. Load a .txt or .pdf file using LangChain's TextLoader and split it using RecursiveCharacterTextSplitter. Print the total number of document chunks created.



```
8 docs.append({"title": os.path.basename(path), "text": text})
9
10 print(f"Loaded {len(docs)} documents from {DOCS_PATH}")
11 for d in docs[:5]:
12     print("-", d["title"], "(", len(d["text"]), "chars)")
13
Loaded 5 documents from /content/my_docs
- refund policy.pdf ( 1712 chars )
- Python_reference.pdf ( 54168 chars )
- cloud computing training.pdf ( 25572 chars )
- Python_reference.pdf ( 54168 chars )
- cloud computing training.pdf ( 25572 chars )

[25] 1 metas = []
2 texts = []
3 for doc in docs:
4     chunks = chunk_text(doc["text"], chunk_size=800, overlap=200)
5     for i, c in enumerate(chunks):
6         metas.append({"title": doc["title"], "chunk_id": i, "text": c})
7         texts.append(c)
8
9 print("Total chunks:", len(texts))
10 if len(texts) == 0:
11     print("No chunks found - check DOCS_PATH, file formats, or chunker settings.")
12
Total chunks: 271

[26] 1 print("Loading embedding model...")
2 embed_model = SentenceTransformer("all-MiniLM-L6-v2")
```

colab.research.google.com/drive/1D0w8XuPsC8Xeg7E7BZ5gqPpY8GVw-gjU#scrollTo=G62F4WjpxHqv

RAG.ipynb

File Edit View Insert Runtime Tools Help

Q Commands + Code + Text ▶ Run all

✓ RAM Disk

260 [22]

```
3 import os
4
5 print("Please upload your ZIP file containing documents (.txt, .md, .pdf)")
6 uploaded = files.upload()
7
8 # Get uploaded file name
9 uploaded_zip = list(uploaded.keys())[0] # First uploaded file
10 target_dir = "/content/my_docs"
11 os.makedirs(target_dir, exist_ok=True)
12
13 # Extract ZIP
14 with zipfile.ZipFile(uploaded_zip, 'r') as zip_ref:
15     zip_ref.extractall(target_dir)
16
17 print(f"Files extracted to: {target_dir}")
18 !ls -R {target_dir}
19
```

Please upload your ZIP file containing documents (.txt, .md, .pdf)

Choose Files data.zip

- data.zip(application/zip) - 1111901 bytes, last modified: 8/11/2025 - 100% done

Saving data.zip to data (1).zip

Files extracted to: /content/my_docs

/content/my_docs:

data myfolder

/content/my_docs/data:

'cloud computing training.pdf' Python_reference.pdf

/content/my_docs/myfolder:

'cloud computing training.pdf' Python_reference.pdf 'refund policy.pdf'

Variables Terminal

12:10 PM T4 (Python 3)

Type here to search

28°C Light rain 12:12 PM 8/11/2025