Topic 2: Open Source LLMs & Local Setup

Assignment 1: Hugging Face Model Exploration

- **Objective:** Explore an open-source LLM on Hugging Face Hub.
- **Instructions:**
    1. Create a Python virtual environment and install `transformers` and `torch`.
    2. Pick a small open-source model (like `distilbert-base-uncased`).
    3. Load the model and tokenizer in Python.
    4. Use it to perform a simple task (e.g., text classification, summarization).
- **Deliverables:** Python script with the following:
    - Environment setup commands
    - Model loading code
    - Sample output from test text

Environmental Setup commands:

In Powershell:

Step 1: Open PowerShell as Administrator

- Click **Start Menu → type "PowerShell" → Right click → Run as Administrator**

Step 2: Change Execution Policy (Safe Way)

Run this command:

```
Set-ExecutionPolicy RemoteSigned -Scope CurrentUser
```

- RemoteSigned = lets you run **local scripts** (like activate.ps1) without restrictions, but still protects against untrusted scripts from the internet.
- -Scope CurrentUser = only applies to **your user account** (not system-wide, safe to use).

Press **Y** (Yes) when prompted.

Step 3: Verify the Change

Run:

```
Get-ExecutionPolicy -Scope CurrentUser
```

It should show:

RemoteSigned

Step 4: Restart VS Code

Now, open VS Code normally. When you create or activate a virtual environment, just do:

.\hf_env\Scripts\activate

Inside VS Code

1. Create Virtual Environment

In VS Code terminal (PowerShell or CMD), run:

```
python -m venv hf_env
```

2. Activate Environment

On Windows:

```
.\hf_env\Scripts\activate
```

3. Install Required Libraries

```
pip install torch transformers
```

```
from transformers import AutoTokenizer, AutoModelForSequenceClassification, pipeline
```

4. Run Your Python Script

- Create a file, e.g., `huggingface.py`
- Add Hugging Face code (sentiment, summarization, classification).
- Run it in terminal:

```
python huggingface.py
```

Output Screenshots:

1. Sentimental analysis
   Model used : distilbert-base-uncased-finetuned-sst-2-english

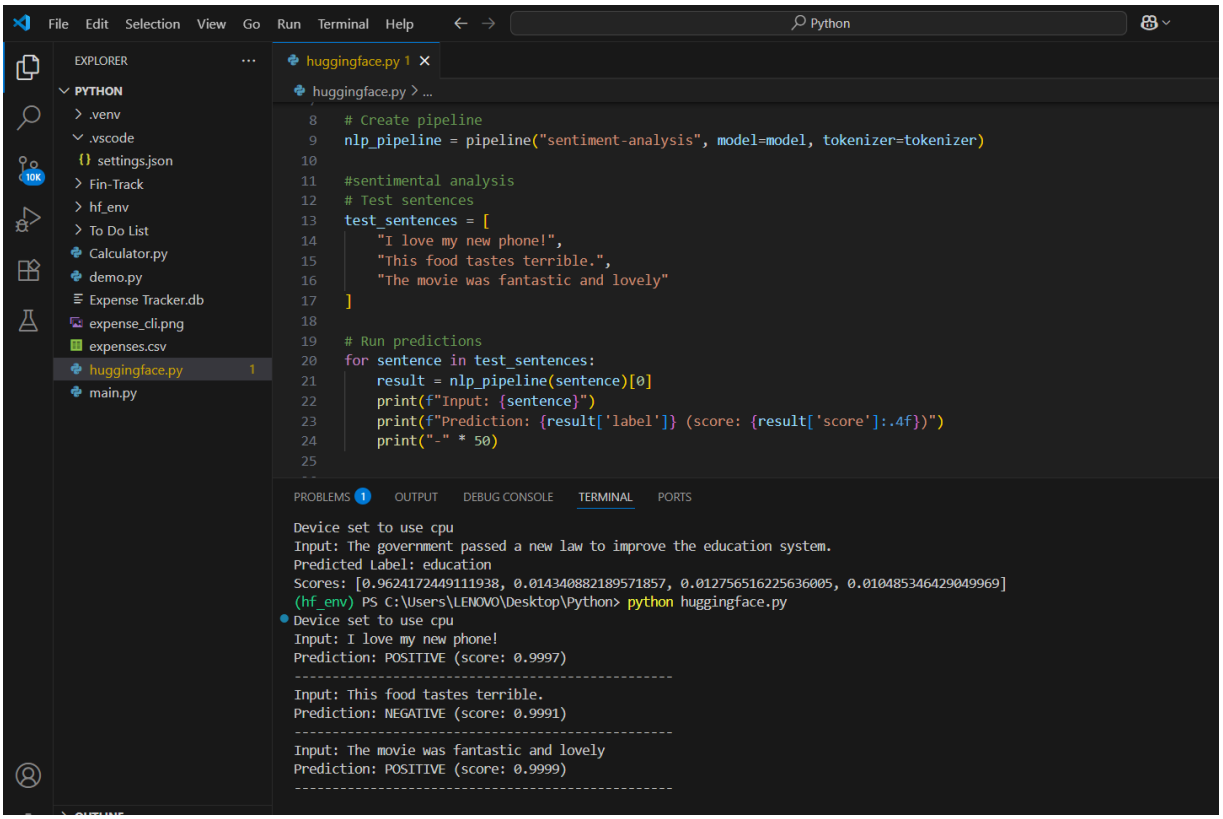

2. Text Summarization
   Model used: sshleifer/distilbart-cnn-6-6

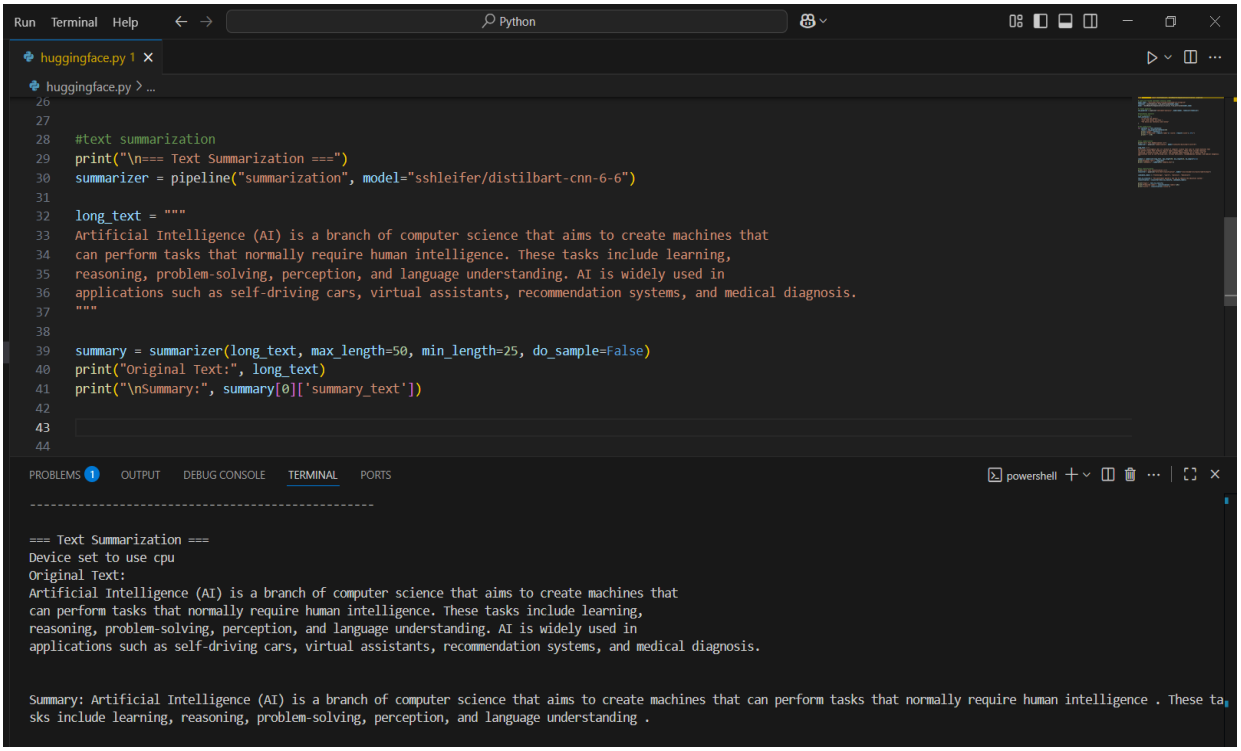

3. Text Classification
   Model used: cross-encoder/nli-distilroberta-base

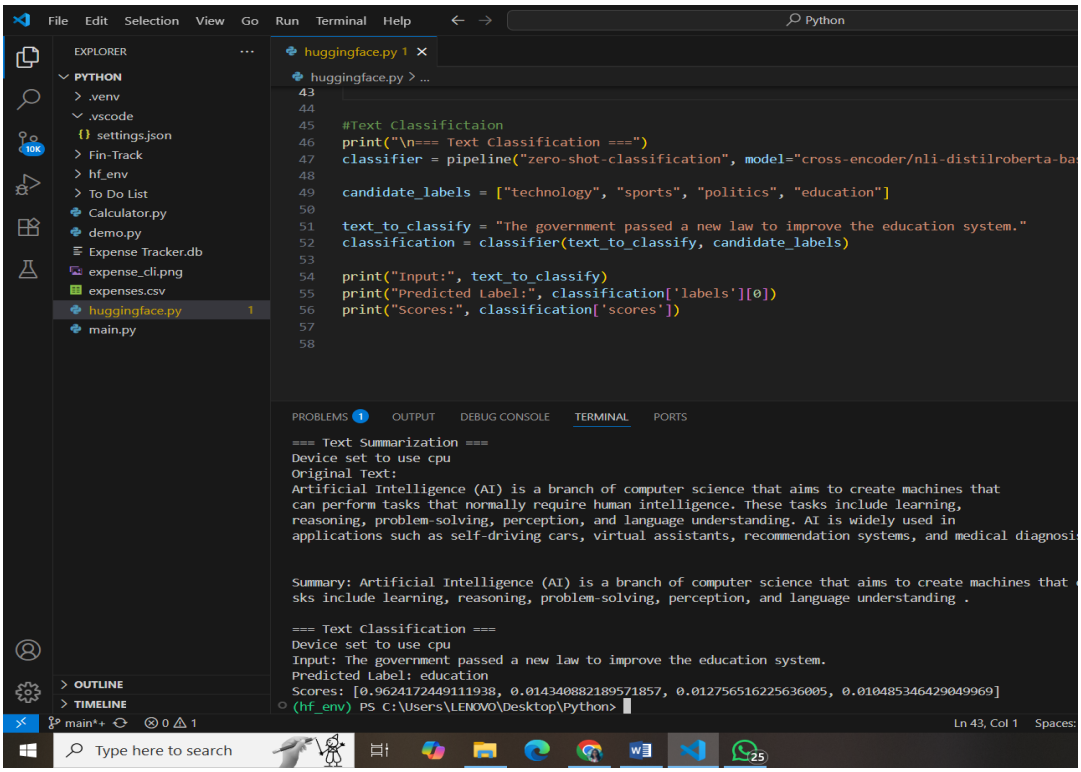

Assignment 2: Local LLM Installation and Testing

- **Objective:** Install a local LLM and test its functionality.
- **Instructions:**
    1. Install **Ollama** or any local LLM (e.g., LLaMA 3.2).
    2. Run a simple prompt like: “Write a short poem about AI.”
    3. Measure the response time and note any errors.
    4. Document any troubleshooting steps you had to take (if installation failed or model crashed).
- **Deliverables:**
    - Screenshot of working LLM
    - Sample prompt and output
    - Short reflection on performance and installation experience

Prompts and output:

On cmd:

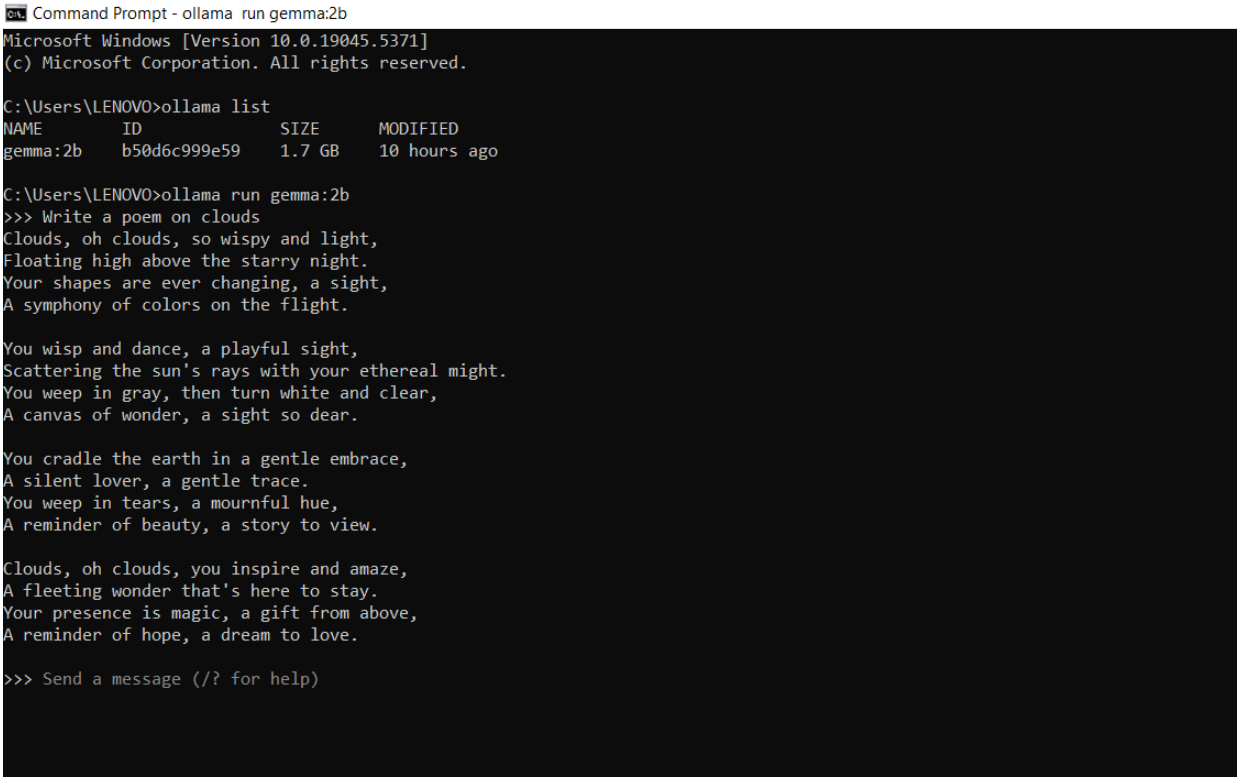

On ollama App:

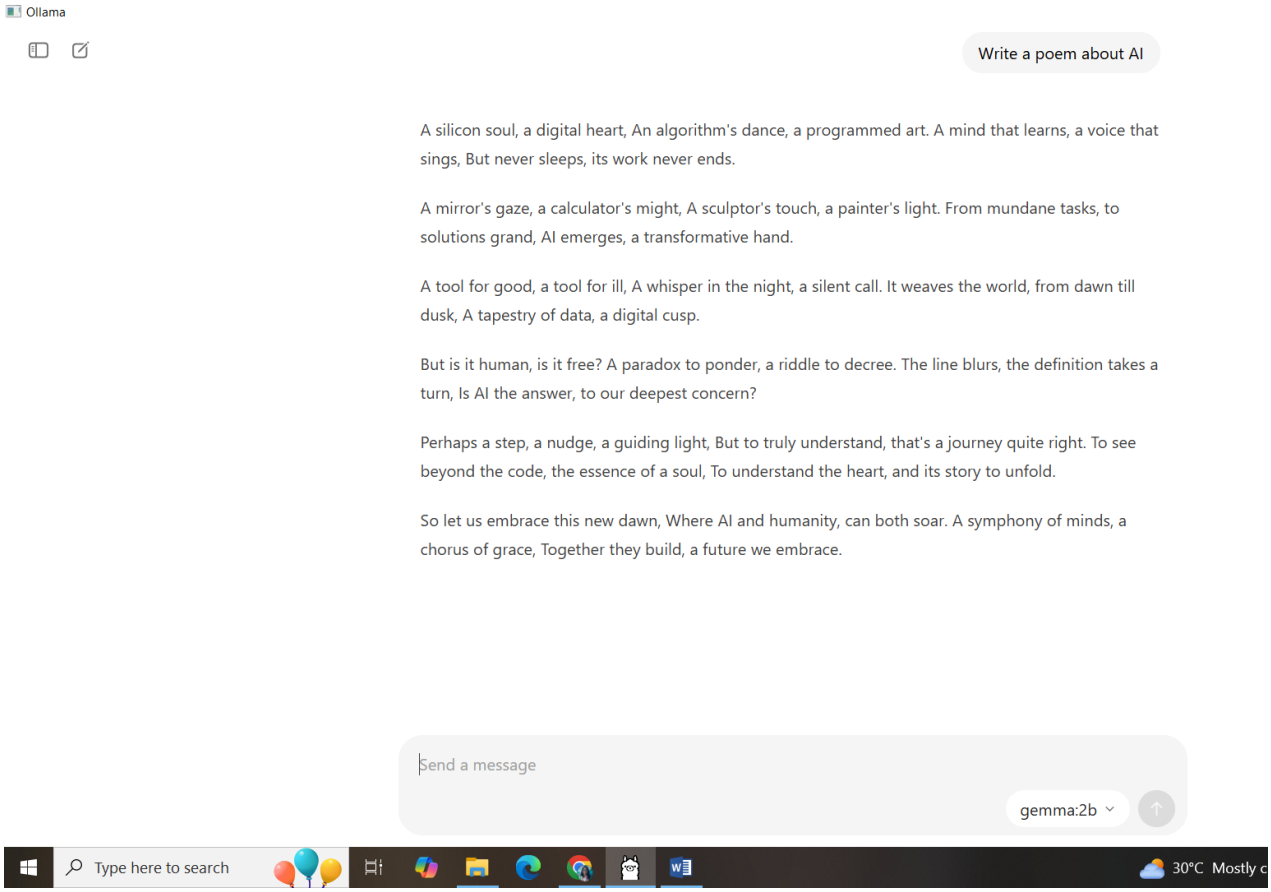

Reflection:

Installation Experience:

The installation of Ollama was quite straightforward. I simply downloaded it from the official Ollama website and verified the installation through the command line. Using commands like ollama list, I could check the available models on my local device. Since this was a fresh installation, no models were present initially, so I decided to add **gemma:2b**, which is around 1.7GB in size. I chose this model because it is relatively lightweight yet effective, while larger models such as **llama2** or **Facebook BART** were much heavier and would have taken considerably more time and space to download. After setting up gemma:2b, I was able to test it both in the command line and in the Ollama app. Overall, the setup was smooth and easy to follow.

Performance:

In terms of performance, downloading the **gemma:2b** model took about **5–10 minutes**, which was reasonable compared to the larger models that required significantly more time. Once downloaded, the model responded well to prompts. When I entered a query, it initially took **10–15 seconds** to process, and then the text started generating at a steady pace, roughly a word per second. The output quality was quite impressive—the responses were coherent and creative, especially when generating a poem. I was satisfied with both the speed and the quality of the results, making gemma:2b a practical choice for my local testing.